

ToneCanvas: Visually Painting Character Tone for Narrative Consistency

Junghwan Kim*
Department of Intelligence
and Information
Seoul National University
Seoul, Republic of Korea
jhbale11@snu.ac.kr

Hana Oh*
Department of Intelligence
and Information
Seoul National University
Seoul, Republic of Korea
hana2001@snu.ac.kr

Bongwon Suh
Department of Intelligence
and Information
Seoul National University
Seoul, Republic of Korea
bongwon@snu.ac.kr




Figure 1: ToneCanvas supports character-tone revision through comparison, specification, and reference. (A) Persistent watercolor overlays make character tone visible directly on the manuscript, while painting-inspired tools support tone sampling, revision, and terminology review. (B) The Tone Palette lets writers inspect and adjust tone through independent dimensions with a live preview. (C) Character Layers gather dispersed dialogue, character references, and reusable Tone Chips to support cross-passage revision.

Abstract

Revising long-form fiction requires writers to maintain each character’s voice across dialogue that is often dispersed throughout a manuscript. Because tone remains embedded in plain text, identifying gradual shifts in how a character speaks relies heavily on rereading, memory, and external references. We present ToneCanvas, an interactive writing interface that supports character-centered tone revision through persistent visual comparison within and across passages. A formative study with five experienced creative writing professionals revealed four recurring breakdowns in character-tone revision, which informed the design of ToneCanvas through persistent visual overlays, character-focused dialogue views, and

adjustable tone dimensions. ToneCanvas contributes a character-centered interface for inspecting and revising character tone within and across dispersed dialogue.

Keywords

Narrative Consistency; Character Voice; Visual Text Analytics; LLM-Augmented Interface

ACM Reference Format:

Junghwan Kim, Hana Oh, and Bongwon Suh. 2026. ToneCanvas: Visually Painting Character Tone for Narrative Consistency. In *The 39th Annual ACM Symposium on User Interface Software and Technology (UIST Adjunct '26)*, November 02–05, 2026, Detroit, MI, USA. ACM, New York, NY, USA, 3 pages. <https://doi.org/10.1145/3830397.3841808>

1 Introduction and Background

Revising long-form fiction requires writers to maintain each character’s voice across dialogue that is often scattered throughout the manuscript. As stories grow longer, keeping track of these voices becomes increasingly difficult. For example, a writer may decide that a villain should gradually become warmer after a narrative turning point, yet identifying every dialogue passage that should reflect this change requires substantial rereading, cross-referencing

*Both authors contributed equally to this research.



with external character notes, and relying on memory across the manuscript [4, 5].

ToneCanvas addresses this challenge by making one important aspect of character voice—tone—visible during revision. We define tone as a character’s habitual manner of speaking, operationalized through dimensions grounded in sociolinguistic research: register and formality [1, 15], politeness [2, 3], directness and intensity [13], and interpersonal stance such as warmth [8]. Unlike emotion, which may fluctuate from scene to scene, tone often reflects a character’s recognizable speaking style while gradually evolving throughout the narrative. While dimensions such as warmth may appear emotion-like, ToneCanvas treats them as properties of how an utterance is delivered rather than what the character feels: a consistently cold character may experience joy in a scene without her dialogue warming, and it is this delivery that the system surfaces and adjusts. Revising tone is therefore an inherently comparative task: writers must inspect how the same character speaks across dispersed dialogue and determine whether observed changes reflect intentional character development or unintended tonal drift.

Existing writing tools provide limited support for reviewing character tone across a manuscript. LLM writing assistants can help with drafting, rewriting, and character development, but they do not relate local revisions to a character’s voice across dispersed dialogue [9, 10, 14, 19]. Recent HCI systems use visual interfaces to support writing by visualizing feedback, story entities and timelines, and LLM-derived narrative trajectories [12, 18, 20]. Narrative visualization systems similarly visualize story structure, character emotion, and character relationships, but stop short of supporting revision of how a character speaks across a manuscript [6, 7, 17]. Drawing-inspired text editing further demonstrates how visual metaphors can support direct manipulation of stylistic properties [11]. ToneCanvas builds on these directions while focusing on character-centered, cross-passage tone revision.

We present ToneCanvas, which reimagines character-tone revision through persistent visual comparison within and across passages. By externalizing character tone onto the manuscript, ToneCanvas helps writers compare dispersed dialogue, inspect possible tonal drift, and revise tone through structured visual interactions. We contribute (1) a formative study identifying four recurring breakdowns in character-tone revision, (2) four design opportunities for supporting distributed, character-centered tone review, and (3) an interactive prototype that realizes these opportunities through persistent visual comparison and parameterized tone control.

2 Formative Study

We interviewed five creative writing professionals (5–15 years of experience) who regularly revise multi-character fiction. The study explored how writers revise long-form fiction to identify recurring revision breakdowns and opportunities for visual support. Across the interviews, maintaining character voice consistently emerged as a shared challenge, with character tone repeatedly identified as difficult to locate, compare, and revise. The 60-minute semi-structured sessions combined workflow walkthroughs, recall of difficult revision episodes, and critiques of early interface sketches. The interviews revealed four recurring breakdowns.

Tone is distributed and hard to surface. Writers relied on rereading and memory to compare how a character spoke across dispersed dialogue. F3, who writes long-running martial-arts web novels, said: “I can’t re-read everything. But there’s no other way.” Gradual shifts were hardest: F5 recalled a villain whose speech “gradually softened over twenty chapters” before a reader noticed.

Writers sense that something is off but cannot specify it. Even after locating a problematic passage, writers often struggled to identify which aspect of the character’s tone had changed. F4 captured the gap between intuition and action: “I can tell something’s off, but then I’m stuck thinking—what should I change it to?” Prompts such as “make this sound colder” often shifted formality, warmth, and directness together, even when only one was intended.

Character voice is documented outside the manuscript. Writers maintained character sheets, glossaries, and reference passages describing how each character should speak. F2 recorded “vocabulary level, sentence length, how often they show emotion, and signature endings” for major characters. Consulting these references required repeatedly switching between the manuscript and supporting documents; participants described this switching as disruptive to revision itself, as each lookup forced them to re-read the passage to re-establish where they were and what they had intended to change.

Names, aliases, and terms vary across long drafts. Writers also struggled to maintain consistent references to characters and locations across long manuscripts. A character might appear under a full name, shortened form, title, alias, or romanization variant. F3 kept a spreadsheet to track these references, but noted “going back and forth is painful,” pointing to a need to review variants directly within the manuscript. Such referential variation is itself part of voice: shifts among a full name, nickname, title, or honorific signal changes in register and interpersonal stance [2], so uncontrolled variation can both obscure the tonal signal writers are tracking and read as unintended tonal drift.

3 ToneCanvas Prototype

Character-tone revision requires comparing how a character speaks across dispersed dialogue. To support this process, ToneCanvas makes character tone persistently visible through a watercolor-inspired visual layer. Watercolor was chosen because it can remain continuously visible without overwhelming the manuscript, allowing tone to be compared during ordinary reading rather than only on demand. The prototype organizes revision around three complementary activities: comparing character tone, specifying tonal changes, and referencing character information.

Compare character tone across passages. ToneCanvas renders dialogue blocks with persistent watercolor overlays that encode prominent tonal dimensions. Color differences draw attention to passages whose tone departs from surrounding dialogue, while  **Focus Mode** gathers one character’s dialogue into a single view for inspecting gradual shifts and recurring voice patterns.

Specify and revise tone. Following structured LLM interfaces that expose an editable design space [16, 21], the Tone Palette lets writers inspect and adjust tone through independently controllable dimensions—separating aspects such as formality, warmth, and directness that writers in our formative study reported shifting

together unintentionally—while previewing changes on a character-specific reference sentence. ✍ **Spoid** samples a passage’s tonal profile into the palette; 🖌 **Brush** applies a target tone with an inline diff for writer approval; 🧽 **Sponge** explores smoother transitions between adjacent passages; and 🎛 **Tone Chips** save reusable tone presets for each character.

Reference character information during revision. Each Character Layer stores a Character Brief containing voice notes, addressing style, and other character-specific references beside the manuscript. These references are automatically incorporated into rewriting so writers do not need to repeatedly consult external documents. 📌 **Stamp** groups names, aliases, titles, and recurring terms—which themselves carry register and interpersonal stance—allowing writers to review and resolve referential variants directly within the manuscript.

4 Conclusion and Future Work

ToneCanvas reframes character-tone revision through persistent visual comparison, making character tone visible within and across dispersed dialogue. Grounded in a formative study with experienced fiction writers and previous literature, the prototype supports comparing character tone, specifying tonal changes, and referencing character information during revision. Future work will evaluate ToneCanvas in writers’ own drafts to understand how persistent visual comparison supports intentional character development over extended writing sessions. We also plan to explore macro-level overviews that summarize tonal trajectories per scene or act, and addressee-dependent tone, as a character’s voice may legitimately shift with the interlocutor.

References

- [1] Douglas Biber. 2006. University language: A corpus-based study of spoken and written registers.
- [2] Penelope Brown. 1987. *Politeness: Some universals in language usage*. Vol. 4. Cambridge university press.
- [3] Cristian Danescu-Niculescu-Mizil, Moritz Sudhof, Dan Jurafsky, Jure Leskovec, and Christopher Potts. 2013. A computational approach to politeness with application to social factors. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Hinrich Schuetze, Pascale Fung, and Massimo Poesio (Eds.). Association for Computational Linguistics, Sofia, Bulgaria, 250–259. <https://aclanthology.org/P13-1025/>
- [4] David K Farkas. 1985. The concept of consistency in writing and editing. *Journal of Technical Writing and Communication* 15, 4 (1985), 353–364. doi:10.2190/T6EM-UTTO-EL6J-59N9
- [5] Linda Flower, John R Hayes, Linda Carey, Karen Schriver, and James Stratman. 1986. Detection, diagnosis, and the strategies of revision. *College Composition & Communication* 37, 1, 16–55. doi:10.58680/ccc198611245
- [6] Md Naimul Hoque, Bhavya Ghai, and Niklas Elmqvist. 2022. DramatVis Personae: Visual Text Analytics for Identifying Social Biases in Creative Writing. In *Proceedings of the 2022 ACM Designing Interactive Systems Conference (Virtual Event, Australia) (DIS '22)*. Association for Computing Machinery, New York, NY, USA, 1260–1276. doi:10.1145/3532106.3533526
- [7] Md Naimul Hoque, Bhavya Ghai, Kari Kraus, and Niklas Elmqvist. 2023. Portrayal: Leveraging NLP and Visualization for Analyzing Fictional Characters. In *Proceedings of the 2023 ACM Designing Interactive Systems Conference (Pittsburgh, PA, USA) (DIS '23)*. Association for Computing Machinery, New York, NY, USA, 74–94. doi:10.1145/3563657.3596000
- [8] Ken Hyland. 2005. Stance and engagement: A model of interaction in academic discourse. *Discourse studies* 7, 2 (2005), 173–192. doi:10.1177/1461445605050365
- [9] Mina Lee, Katy Ilonka Gero, John Joon Young Chung, Simon Buckingham Shum, Vipul Raheja, Hua Shen, Subhashini Venugopalan, Thiemo Wambsgans, David Zhou, Emad A. Alghamdi, Tal August, Avinash Bhat, Madiha Zahrah Choksi, Senjuti Dutta, Jin L.C. Guo, Md Naimul Hoque, Yewon Kim, Simon Knight, Seyed Parsa Neshaei, Antonette Shibani, Disha Shrivastava, Lila Shroff, Agnia Sergeyuk, Jessi Stark, Sarah Sberman, Sitong Wang, Antoine Bosselut, Daniel Buschek, Joseph Chee Chang, Sherol Chen, Max Kreminski, Joonsuk Park, Roy Pea, Eugenia Ha Rim Rho, Zejiang Shen, and Pao Siangliulue. 2024. A Design Space for Intelligent and Interactive Writing Assistants. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '24). Association for Computing Machinery, New York, NY, USA, Article 1054, 35 pages. doi:10.1145/3613904.3642697
- [10] Mina Lee, Percy Liang, and Qian Yang. 2022. CoAuthor: Designing a Human-AI Collaborative Writing Dataset for Exploring Language Model Capabilities. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems* (New Orleans, LA, USA) (CHI '22). Association for Computing Machinery, New York, NY, USA, Article 388, 19 pages. doi:10.1145/3491102.3502030
- [11] Damien Masson, Young-Ho Kim, and Fanny Chevalier. 2025. Textshop: Interactions Inspired by Drawing Software to Facilitate Text Editing. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, Article 1087, 14 pages. doi:10.1145/3706598.3713862
- [12] Damien Masson, Zixin Zhao, and Fanny Chevalier. 2025. Visual Story-Writing: Writing by Manipulating Visual Representations of Stories. In *Proceedings of the 38th Annual ACM Symposium on User Interface Software and Technology (UIST '25)*. Association for Computing Machinery, New York, NY, USA, Article 70, 15 pages. doi:10.1145/3746059.3747758
- [13] Saif Mohammad. 2018. Obtaining Reliable Human Ratings of Valence, Arousal, and Dominance for 20,000 English Words. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Iryna Gurevych and Yusuke Miyao (Eds.). Association for Computational Linguistics, Melbourne, Australia, 174–184. doi:10.18653/v1/P18-1017
- [14] Hua Xuan Qin, Shan Jin, Ze Gao, Mingming Fan, and Pan Hui. 2024. CharacterMeet: Supporting Creative Writers’ Entire Story Character Construction Processes Through Conversation with LLM-Powered Chatbot Avatars. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '24). Association for Computing Machinery, New York, NY, USA, Article 1051, 19 pages. doi:10.1145/3613904.3642105
- [15] Sudha Rao and Joel Tetreault. 2018. Dear Sir or Madam, May I Introduce the GYAFC Dataset: Corpus, Benchmarks and Metrics for Formality Style Transfer. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, Marilyn Walker, Heng Ji, and Amanda Stent (Eds.). Association for Computational Linguistics, New Orleans, Louisiana, 129–140. doi:10.18653/v1/N18-1012
- [16] Sangho Suh, Meng Chen, Bryan Min, Toby Jia-Jun Li, and Haijun Xia. 2024. Luminate: Structured Generation and Exploration of Design Space with Large Language Models for Human-AI Co-Creation. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '24). Association for Computing Machinery, New York, NY, USA, Article 644, 26 pages. doi:10.1145/3613904.3642400
- [17] Katie Watson, Samuel S. Sohn, Sasha Schriber, Markus Gross, Carlos Manuel Muniz, and Mubbasir Kapadia. 2019. StoryPrint: an interactive visualization of stories. In *Proceedings of the 24th International Conference on Intelligent User Interfaces (Marina del Ray, California) (IUI '19)*. Association for Computing Machinery, New York, NY, USA, 303–311. doi:10.1145/3301275.3302302
- [18] Catherine Yeh, Tara Menon, Robin Singh Arya, Helen He, Moira Weigel, Fernanda Viégas, and Martin Wattenberg. 2025. Story Ribbons: Reimagining Storyline Visualizations with Large Language Models. *IEEE Transactions on Visualization and Computer Graphics* (2025). doi:10.1109/TVCG.2025.3634265
- [19] Ann Yuan, Andy Coenen, Emily Reif, and Daphne Ippolito. 2022. Wordcraft: Story Writing With Large Language Models. In *Proceedings of the 27th International Conference on Intelligent User Interfaces (Helsinki, Finland) (IUI '22)*. Association for Computing Machinery, New York, NY, USA, 841–852. doi:10.1145/3490099.3511105
- [20] Chao Zhang, Kexin Ju, Zhuolun Han, Yu-Chun Grace Yen, and Jeffrey M. Rzeszutski. 2025. Synthia: Visually Interpreting and Synthesizing Feedback for Writing Revision. In *Proceedings of the 38th Annual ACM Symposium on User Interface Software and Technology (UIST '25)*. Association for Computing Machinery, New York, NY, USA, Article 88, 16 pages. doi:10.1145/3746059.3747703
- [21] Zheng Zhang, Jie Gao, Ranjodh Singh Dhaliwal, and Toby Jia-Jun Li. 2023. VISAR: A Human-AI Argumentative Writing Assistant with Visual Programming and Rapid Draft Prototyping. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (San Francisco, CA, USA) (UIST '23)*. Association for Computing Machinery, New York, NY, USA, Article 5, 30 pages. doi:10.1145/3586183.3606800